

An Introduction to Regularized Regression

Machine Learning and Causal
Inference

Susan Athey

Thanks to Sendhil Mullainathan for sharing slides; see Mullainathan and Speiss (2017) JEP which covers much of this material

What we do in Econometrics: The Case of Regression

- Specify a model:

$$Y_i = f(X_i) + \epsilon_i = X_i\beta + \epsilon_i$$

- Data set has observations $i=1, \dots, n$
- Use OLS regression on the entire dataset to construct an estimate $\hat{\beta}$
- Discuss assumptions under which some components of $\hat{\beta}$ have a causal interpretation
- Consider that S_n (set of observed units, $i=1, \dots, n$) is a random sample from a much larger population.
- Construct confidence intervals and test the hypothesis that some components are equal to zero.
- Theorem: OLS is BLUE (Best Linear Unbiased Estimator)
 - Best = lowest-variance

Goals of Prediction and Estimation

- Goal of estimation: unbiasedness

$$E[\hat{f}] = f$$

- Goal of prediction: loss minimization

$$L(f) = E_{(x,y)} \ell(f(x), y)$$

$$\hat{f} \approx \min_{f \in \mathcal{F}} L(f)$$

- E.g. $\ell(f(x), y) = (f(x) - y)^2$
- Use the data to pick a function that does well on a new data point

Key assumptions in both cases

- Stationary data generating process

Data

$$S_n = (y_i, x_i) \text{ iid}$$

- Estimation:
 - Interested in a parameter of that process
- Prediction:
 - Interested in predicting y

High v. Low Dimensional Analysis

- We have discussed prediction as a high dimensional construct
- Practically that is where it is useful
- But to understand how high dimensional prediction works we must unpack an implicit presumption
 - Presumption: Our known estimation strategies would be great predictors *if they were feasible*

A Simple OLS example

- Suppose we truly live in a linear world

$$y = \beta_0 + \beta_1 x_1 + \varepsilon$$

$$\varepsilon \sim \mathcal{N}(0, \sigma_\varepsilon^2) \quad x_1 \sim N(0, 1)$$

- Write $x = (1, x)$

$$y = \beta x + \varepsilon.$$

OLS seems like a good predictor

$$L(\hat{f}^{\text{OLS}}) = \mathbb{E}_{(y,x)}(\hat{\beta}'x - y)^2 = (\hat{\beta}_0 - \beta_0)^2 + (\hat{\beta}_1 - \beta_1)^2 + \sigma_\varepsilon^2$$

So wouldn't we want the $\hat{\beta}$ with $\mathbb{E}_{S_n}(\hat{\beta}) = \beta$?

Especially since it is known to be efficient

An Even Simpler Set-up

- Let's get even lower dimensional
- No variables at all
- Suppose you get the data of the type:

$$y_i = \mu + \epsilon_i$$

- You would like to estimate the mean

Forming an estimator of the mean

- Minimize bias: $\alpha = 1$
- The sample mean is an unbiased estimator
 - Also what you would get from OLS regression on a constant


A prediction problem

- In the same setup, you are given n data points
- You would like to guess the value of a new data point from the same distribution
- Goal: minimize quadratic loss of prediction

Best Predictor

$$\ell(\hat{\mu}, y)]$$

The higher
alpha the lower
the bias



The higher alpha the more variable across samples it is

Key problem

- The unbiased estimator has a nice property:

$$E[\hat{\mu}|\mu] = \mu$$

- But getting that property means large sample to sample variation of estimator
- This sample to sample variation means that in any particular finite sample I'm paying the cost of being off on all my predictions

Intuition

- I see your first test score. What should my prediction of your next test be?
 - Your first test score is an unbiased estimator
 - But it is very variable
- Note: “Bayesian” intuition
 - Even simpler: what was my guess before I saw any information
 - Shrink to that
 - In this example I’m shrinking to zero

But in a way you know this

- As empiricists you already have this intuition

Back to Simple OLS example

- Suppose we truly live in a linear world

$$y = \beta_0 + \beta_1 x_1 + \varepsilon$$

$$\varepsilon \sim \mathcal{N}(0, \sigma_\varepsilon^2) \quad x_1 \sim N(0, 1)$$

- Write $x = (1, x)$

$$y = \beta x + \varepsilon.$$

A Simple Example

- You run a one variable regression and get

$$\hat{\beta}_0^{\text{OLS}} = 0 \pm .2$$

$$\hat{\beta}_1^{\text{OLS}} = 2 \pm 10$$

- Would you use the OLS coefficients to predict
- Or drop the first variable and use this:

$$\hat{\beta}_0^{\text{OLS}} = \arg \min_{\beta_0} \widehat{\mathbb{E}}_{S^n} (\beta_0 - y)^2 = \widehat{\mathbb{E}}_{S^n} y$$

Deciding whether to drop

- Suppose in the (impossible) case we got the true world right.
 - $(0,2)$ are the right coefficients
- Of course OLS does perfectly (by assumption).
- But how would OLS do on new samples...where $(0,2)$ being the generating coefficients?
 - We're giving OLS a huge leg up here.

OLS Performance

$$\begin{aligned}
 \mathcal{L}_n(\text{OLS}) - \sigma_\varepsilon^2 &= \\
 &= E_{(y,x)} E_{S_n} [\beta'x - (\hat{\beta}^{\text{OLS}})'x]^2 \\
 &= E_{(y,x)} [\underbrace{(\beta'x - (E_{S_n} \hat{\beta}^{\text{OLS}})'x)^2}_{\text{unbiased}} + \text{Var}_{S_n}((\hat{\beta}^{\text{OLS}})'x)] \\
 &= \text{Var}_{S_n}(\hat{\beta}_0^{\text{OLS}}) + \text{Var}_{S_n}(\hat{\beta}_1^{\text{OLS}})
 \end{aligned}$$

What if we dropped the variable

$$\mathcal{L}_n(\text{OLS}_0) - \sigma_\varepsilon^2 =$$

$$\begin{aligned}
\mathcal{L}_n(\text{OLS}) - \sigma_\varepsilon^2 &= \\
&= E_{(y,x)} E_{S_n} [\beta'x - (\hat{\beta}^{\text{OLS}})'x]^2 \\
&= E_{(y,x)} \left[\underbrace{(\beta'x - (E_{S_n} \hat{\beta}^{\text{OLS}})'x)^2}_{\text{unbiased}} + \text{Var}_{S_n}((\hat{\beta}^{\text{OLS}})'x) \right] \\
&= \text{Var}_{S_n}(\hat{\beta}_0^{\text{OLS}}) + \text{Var}_{S_n}(\hat{\beta}_1^{\text{OLS}})
\end{aligned}$$

$$\begin{aligned}
\mathcal{L}_n(\text{OLS}_0) - \sigma_\varepsilon^2 &= \\
&= E_{(y,x)} E_{S_n} [\beta'x - (\hat{\beta}^{\text{OLS}_0})'x]^2 \\
&= E_{(y,x)} \left[\underbrace{(\beta'x - (E_{S_n} \hat{\beta}^{\text{OLS}_0})'x)^2}_{\text{bias}} + \text{Var}_{S_n}((\hat{\beta}^{\text{OLS}_0})'x) \right] \\
&= (0 - 2)^2 + \text{Var}_{S_n}(\hat{\beta}_0^{\text{OLS}_0}) + \text{Var}_{S_n}(\hat{\beta}_1^{\text{OLS}_0})
\end{aligned}$$


$$\mathcal{L}_n(\text{OLS}) - \mathcal{L}_n(\text{OLS}_0) = \underbrace{\text{Var}(\hat{\beta}_1^{\text{OLS}})}_{\text{variance}} - \underbrace{(0 - 2)^2}_{\text{bias}}$$

↑
Your standard error worry!


Where does your standard error intuition come from?

- You see a standard error
- You think “that variable is not ‘significant’” so you might not want to include it.
- But this is misleading


$\hat{\beta}$



$\hat{\beta}$




$\hat{\beta}$



$\hat{\beta}$

Variability
of Estimator

Id have
been generated?



Your Standard Error Worry

- For hypothesis testing se tells you whether the coefficient is significant or not
- For prediction it's telling you how variable an estimator using it really is

Dual purposes of the standard error

- The standard error also tells you that even if you're right on average:
 - Your estimator will produce a lot of variance
 - And then in those cases you make systematic prediction mistakes.
- Bias variance tradeoff
 - Being right on average on the coefficient is not equal to the best predictor.

The Problem Here

- Prediction quality suffers from:
 - Biased coefficients
 - Variability in *estimated* coefficients
 - Even if the true coefficient is 2, in any sample, we will estimate something else
- OLS is lexicographic
 - First ensure unbiased
 - Amongst unbiased estimators: seek efficiency
- Good predictions must trade these off

Two Variable Example

- Belaboring the point here...
- Assume now that we have two variables
 - As before, both normally distributed unit variance
- Your estimator produces

$$\hat{\beta}_0^{OLS} = 0 \pm .2$$

$$\hat{\beta}_1^{OLS} = 2 \pm 10$$

What would you do now?

- Logic above suggests you would drop both variables?
- Or keep both variables?
- It really depends on how you feel about the variance (10)?

Calculation

$$\begin{aligned} \mathcal{L}_n(\text{OLS}) - \mathcal{L}_n(\text{OLS}_0) = & \overbrace{\text{Var}(\hat{\beta}_1^{\text{OLS}}) + \text{Var}(\hat{\beta}_2^{\text{OLS}})}^{\text{variance}} - \overbrace{((0-2)^2 + (0-2)^2)}^{\text{bias}} \\ & + \underbrace{2\rho_{12} \text{Cov}(\hat{\beta}_1^{\text{OLS}}, \hat{\beta}_2^{\text{OLS}})}_{\text{covariance variance}} - \underbrace{2\rho_{12}(0-2)^2}_{\text{covariance bias}} \end{aligned}$$

$$\mathcal{L}_n(\text{OLS}) - \mathcal{L}_n(\text{OLS}_0) = \underbrace{\text{Var}(\hat{\beta}_1^{\text{OLS}})}_{\text{variance}} - \underbrace{(0-2)^2}_{\text{bias}}$$

Hidden in Bias-Variance Tradeoff

- Covariance is central
- The standard error on several variables can be large, even though together their effect is highly consistent
- For prediction covariance between x matters

In a way this problem is not important

$$\mathcal{L}_n(\text{OLS}) - \mathcal{L}_n(\text{OLS}_0) = \underbrace{\text{Var}(\hat{\beta}_1^{\text{OLS}})}_{\text{variance}} - \underbrace{(0 - 2)^2}_{\text{bias}}$$

- The variance term diminishes with sample size
 - Prediction-estimation wedge falls off as $\frac{1}{n}$
- But variance term increases with “variables”
 - Prediction-estimation rises with k
- So this is a problem when...
 - Function class high dimensional relative to data $\frac{k}{n}$

What this means practically

- In some cases what you already know (estimation) is perfectly fine for prediction
 - This is why ML textbooks teach OLS, etc.
 - They are perfectly useful for the kinds of prediction problems ML tries to solve *in low dimensional settings*
- But in high dimensional settings...
 - Note: high dimensional does not ONLY mean lots of variables! It can mean rich interactions.

So far...

- All this gives you a flavor of how the prediction task is not mechanically a consequence of the estimation task
- But it doesn't really tell you **how** to predict
 - Bias variance tradeoff is entirely unactionable
 - What's the bias?
 - What's the variance?
 - This is not really a tradeoff you can make
- A different look at the same problem produces a practical insight though

Back to OLS

$$\hat{\beta}^{\text{OLS}} = \arg \min_{\beta} \hat{\mathbb{E}}_{S_n} (\beta' x - y)^2$$

AVERAGES
NOTATION:
 $\hat{\mathbb{E}}_{S_n}$ for
sample ave.
for sample S_n

$$\beta_{\text{prediction}}^* = \arg \min_{\beta} \mathbb{E}_{(y,x)} (\beta' x - y)^2$$

- The real problem here is minimizing the “wrong” thing: In-sample fit vs out-of-sample fit

Overfit problem

- OLS looks good with the sample you have
 - It's the best you can do *on this sample*
- Bias-variance improving predictive power is about improving *out of sample* predictive power
- Problem is OLS by construction overfits
 - We overfit in estimation

This problem is exactly why wide data is troubling

- Similarly think of the wide data case
- Why are we worried about having so many variables?
- We'll fit very well (perfectly if $k > n$) in sample
- But arbitrarily badly out of sample

Understanding overfit

- Let's consider a general class of algorithms

A General Class of Algorithms

- Let $L(f) = \int_{x,y} \ell(f(x), y) dP(x, y)$ for some loss function ℓ (e.g. squared error)
 - Note: L is an unknown function: we don't know P
- Consider algorithms of the form
$$\hat{f}_{A,S_n} = \arg \min_{f \in \mathcal{F}_A} \hat{L}_{S_n}(f)$$
 - \hat{L}_{S_n} is used here as shorthand for sample mean observations in sample S_n of size n
 - OLS is an *empirical loss minimizer*: it minimizes the sample average over observed data of the loss function
- So empirical loss minimization algorithms are defined by the function class they choose from
- For estimation what we typically do...
 - Show that empirical loss minimizers generate unbiasedness

Empirical Loss minimization

- Leads to unbiasedness/consistency
 - Fit the data you have...
 - In a frequentist world “on average” (across all S_n) this will produce the right thing
 - This is usually how we prove consistency/unbiasedness
- Other variants:
 - MLE

Some Notation

- Define

$$f^* = \arg \min_{f \in \mathcal{F}} L(f) \quad \text{The best we can do}$$

$$f_A^* = \arg \min_{f \in \mathcal{F}_A} L(f) \quad \text{The best in the subset of functions that the algorithm looks at}$$

- Recall: L is infeasible b/c we don't know true data-generating process
- Contrast the latter with:

$$\hat{f}_{A, S_n} = \arg \min_{f \in \mathcal{F}_A} \hat{L}_{S_n}(f)$$

What the in-sample loss minimizer actually produces given a sample

Performance of Algorithm

- Performance of a predictor

$$L(\hat{f}_{A,S_n})$$

- Performance of an Algorithm

$$\mathcal{L}_n(A) := E_{S_n} L(\hat{f}_{A,S_n})$$

- Algorithm's expected loss
- (Suppress S_n in some of the notation for estimator)

The performance of A

$$\mathcal{L}_n(A) = \underbrace{L(f^*)}_{\text{irreducible error}} + \underbrace{L(f_A^*) - L(f^*)}_{\text{approximation error}} + \underbrace{\mathbb{E}_{S_n}(L(\hat{f}_A) - L(f_A^*))}_{\text{estimation error}},$$

Understanding estimation error:

$$\mathbb{E}_{S_n}(\hat{L}(\hat{f}_A) - L(f_A^*)) = \mathbb{E}_{S_n}(\hat{L}(\hat{f}_A) - \hat{L}(f_A^*)) + \mathbb{E}_{S_n}(L(\hat{f}_A) - \hat{L}(\hat{f}_A))$$

“Wrong” function
looks good in-sample

Algorithm does
not see this

Basic Tradeoff

- These two terms go hand in hand:

$$\mathcal{L}_n(A) = \underbrace{L(f^*)}_{\text{irreducible error}} + \underbrace{L(f_A^*) - L(f^*)}_{\text{approximation error}} + \underbrace{E_{S_n}(L(\hat{f}_A) - L(f_A^*))}_{\text{estimation error}},$$

$$E_{S_n}(\underbrace{L(\hat{f}_A) - L(f_A^*)}_{\text{out-of-sample, } \geq 0}) = E_{S_n}(\underbrace{\hat{L}(\hat{f}_A) - \hat{L}(f_A^*)}_{\text{in-sample, } \leq 0}) + \underbrace{E_{S_n}(L(\hat{f}_A) - \hat{L}(\hat{f}_A))}_{\text{unseen overfit}}$$

Approximation – Overfit Tradeoff

- If we reduce set of f to reduce possible over-fit:

$$\underbrace{E_{S_n}(L(\hat{f}_A) - \hat{L}(\hat{f}_A))}_{\text{unseen overfit}}$$

- Then we fit fewer “true functions” and drive up

$$\underbrace{L(\hat{f}_A) - L(f^*)}_{\text{approximation error}}$$

- Only way to avoid this is if we knew information about f^* so we could shrink the set

Unobserved overfit

- So the problem of prediction really is managing unobserved overfit

$$\underbrace{L(\hat{f}_A)}_{\text{unobserved out-of-sample}} = \underbrace{\hat{L}(\hat{f}_A)}_{\text{observed in-sample}} + \underbrace{(L(\hat{f}_A) - \hat{L}(\hat{f}_A))}_{\text{unobserved overfit}}$$

- We do well in-sample. But some of that “fit” is overfit.

Return to the original example

OLS

**Greater Chance
To Overfit**

OLS_0

**Less Chance
To Overfit**

- We drove down overfit by doing a constrained optimization

Basic Tradeoff at the Heart of Machine Learning

- Bigger function classes...
 - The more likely we are to get to the truth (less approximation)
 - The more likely we are to overfit
- So we want to not just minimize in-sample error given a class of functions
- We also want to decide on the class of functions
 - More expressive means less approximation error
 - More expressive means more overfit

Let's do the same thing here

Unconstrained

$$\hat{f}_{A,S_n} = \arg \min_{f \in \mathcal{F}_{\mathcal{A}}} \hat{L}_{S_n}(f)$$

But we are worried about $\underbrace{E_{S_n}(L(\hat{f}_A) - \hat{L}(\hat{f}_A))}_{\text{unseen overfit}}$

So why not do this instead?

$$\arg \min_{f \in \mathcal{F}_{\mathcal{A}}} \hat{L}_{S_n}(f)$$

$$\text{s.t. } R(f) \leq c$$

Complexity measure: tendency to overfit

Return to the original example

OLS

OLS_0

Greater Overfit

Less Overfit

Better approximation

Worse approximation

More **Expressive**

$R(f)$ *higher*

Less **Expressive**

$R(f)$ *lower*

- Reduce overfit by approximating worse
- Choose less expressive function class

Constrained minimization

- We could do a constrained minimization
- But notice that this is equivalent to:

$$\hat{f}_{A_\lambda, S_n} = \arg \min_{f \in \mathcal{F}_{\mathcal{A}}} \hat{L}_{S_n}(f) + \underbrace{\lambda R(f)}_{\text{want: } \approx L(f) - \hat{L}(f)}$$

- Complexity measure should capture tendency to overfit

Basic insight

- Data has signal and noise
- More expressive function classes-
 - Allow us to pick up more of the signal
 - But also pick up more of the noise
- So the problem of prediction becomes the problem of *choosing expressiveness*

Overall Structure

- Create a regularizer that:
 - Measures expressiveness
- Penalize algorithm for choosing more expressive functions
 - Tuning parameter λ
- Let it weigh this penalty against in-sample fit

Linear Example

- Linear function class $x \mapsto \beta'x$ ($\beta \in \mathbb{R}^{k+1}$)

- Regularized linear regression

$$\hat{\beta}_{\lambda}^R = \arg \min_{\beta \in \mathbb{R}^{k+1}} \widehat{\mathbb{E}}_{S_n} (\beta'x - y)^2 - \lambda R(\beta)$$

Regularizers for Linear Functions

- Linear functions more expressive if use more variables

$$R(\beta) = \sum_{j=1}^k 1_{\beta_j \neq 0}$$

- Can transform coefficients

$$R(\beta) = \sum_{j=1}^k |\beta_j|^p$$


Computationally More Tractable

- Lasso

$$\mathcal{F}_{1,c} = \{f_{\gamma}; \sum_{j=1}^k |\gamma_j| \leq c\}$$

- Ridge

$$\mathcal{F}_{2,c} = \{f_{\gamma}; \sum_{j=1}^k \gamma_j^2 \leq c\}$$



What makes a good regularizer?

- You might think...
 - Bayesian assumptions
 - Example: Ridge
- A good regularizer can build in beliefs
- Those are great and useful when available
- But central force is tendency to overfit
- Example:
 - Even if true world were not sparse or priors were not normal you'd still do this

Summary

- Regularization is one half of the secret sauce
- Gives a single-dimensional way of deciding of capturing expressiveness

$$\hat{f}_{A_\lambda, S_n} = \arg \min_{f \in \mathcal{F}_{\mathcal{A}}} \hat{L}_{S_n}(f) + \boxed{\lambda} R(f)$$

- Still missing ingredient is lambda

Choosing lambda

- How much should we penalize expressiveness?
- How do you make the over-fit approximation tradeoff?
- The **tuning** problem.
- Use cross-validation

How Does Cross Validation Work?




Tuning Set = $\frac{1}{5}$ of Training Set

Cross-Validation Mechanics

- Loop over cross-validation samples
 - Train a deep tree on CV-training subset
- Loop over penalty parameters λ
 - Loop over cross-validation samples
 - Prune the tree according to penalty
 - Calculate new MSE of tree
 - Average (over c-v samples) the MSE for this penalty
- Choose the penalty λ^* that gives the best average MSE

LASSO c-v Example

```
plot(lasso.linear)
```



Creating Out-of-Sample In Sample

- Major point:
 - Not many assumptions
 - Don't need to know true model.
 - Don't need to know much about algorithm
- Minor but important point
 - To get asymptotics right we need to make some regularity assumptions
- Side point (to which we return)
 - We'd like to choose best algorithm for sample size n
 - But this will not do that. **Why?**

Why does this work?

1. Not just because we can split a sample and call it out of sample
 - It's because the thing we are optimizing is observable (easily estimable)

This is more than a trick

- It illustrates what separates prediction from estimation:
 - I can't 'observe' my prior.
 - Whether the world is truly drawn from a linear model
 - But prediction quality is observable
- Put simply:
 - Validity of predictions are measurable
 - Validity of coefficient estimators require structural knowledge

This is the essential ingredient to prediction: Prediction quality is an empirical quantity not a theoretical guarantee

Why does this work?

1. It's because the thing we are optimizing is observable
2. By focusing on prediction quality we have reduced dimensionality

To understand this...

- Suppose you tried to use this to choose coefficients
 - Ask which set of coefficients worked well out-of sample.
- Does this work?
- Problem 1: Estimation quality is unobservable
 - Need the same assumptions as algorithm to know whether you “work” out of sample
 - If you just go by fit you are ceding to say you want best predicting model
- Problem 2: No dimensionality reduction.
 - You’ve got as many coefficients as before to search over

$$\hat{\beta}_{\lambda}^R = \arg \min_{\beta \in \mathbb{R}^{k+1}} \mathbb{E}_{S_n} (\beta' x - y)^2 + \lambda R(\beta)$$

Method	$R(\beta)$
OLS	0
Subset selection	$\ \beta\ _0 = \sum_{j=1}^k \mathbb{1}_{\beta_j \neq 0}$
Lasso	$\ \beta\ _1 = \sum_{j=1}^k \beta_j $
Ridge	$\ \beta\ _2^2 = \sum_{j=1}^k \beta_j^2$
Elastic Net	$\alpha \ \beta\ _1 + (1 - \alpha) \ \beta\ _2^2$

Bayesian Interpretation of Ridge

Consider the regression

$$Y_i = \sum_{k=1}^K \beta_k \cdot X_{ik} + \varepsilon_i$$

with

$$\varepsilon_i | X_{i1}, \dots, X_{iK} \sim \mathcal{N}(0, \sigma^2)$$

Suppose we put a prior on the β_k :

$$\beta_k \sim \mathcal{N}(0, \tau^2)$$

and all the β_k independent. Assume σ^2 is known.

Bayesian Interpretation of Ridge

Then the posterior distribution is proportional to

$$\begin{aligned} p(\beta|\text{data}) &\propto \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^N \left(Y_i - \sum_{k=1}^K \beta_k \cdot X_{ik}\right)^2\right) \prod_{k=1}^K \exp\left(-\frac{\beta_k^2}{2\tau^2}\right) \\ &= \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^N \left(Y_i - \sum_{k=1}^K \beta_k \cdot X_{ik}\right)^2 - \sum_{k=1}^K \frac{\beta_k^2}{2\tau^2}\right) \\ &= \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^N (Y_i - \beta' X_i)^2 - \frac{\beta' \beta}{2\tau^2}\right) \end{aligned}$$

Bayesian Interpretation of Ridge

So, the posterior is normal, and the posterior mean minimizes

$$\begin{aligned} \sum_{i=1}^N (Y_i - \beta' X_i)^2 + \beta' \beta \cdot \frac{\sigma^2}{\tau^2} \\ = \sum_{i=1}^N (Y_i - \beta' X_i)^2 + \frac{\sigma^2}{\tau^2} \cdot \|\beta\|^2 \end{aligned}$$

This leads to the posterior mean


$$(\mathbf{X}'\mathbf{X} + I_K \cdot \sigma^2/\tau^2)^{-1} \mathbf{X}'\mathbf{Y}.$$

If the $\mathbf{X}'\mathbf{X}$ matrix is diagonal, all elements of β would be shrunk towards zero by the same fraction. With a non-diagonal matrix the degree of shrinkage varies.


POST-Lasso

- Important distinction:
 - Use LASSO to choose variables
 - Use OLS on these variables
- How should we think about these?

Ridge regularization path




Lasso regularization path



In the orthonormal case, i.e. $\mathbf{X}^T \mathbf{X} = \mathbf{I} = (\mathbf{X}^T \mathbf{X})^{-1}$:


$$\hat{\beta}_j(\lambda_1) = \text{sgn}(\hat{\beta}_j) (|\hat{\beta}_j| - \lambda_1/2)_+$$

That is, the lasso estimate is related to the OLS estimate via the so-called *soft threshold function* (depicted here for $\lambda=1$).






Why not Hard Thresholding?

Soft Thresholding



Orthonormal: $\hat{\beta}_{RIDGE} = \frac{\hat{\beta}_{OLS}}{1 + \lambda}$





Can be very misleading

A Basic Policy Problem

- Every transfer program in the world must determine...
 - Who is eligible for the transfer
- Typical goal of redistributive programs
 - Transfer to neediest
- But identifying the neediest is easier said than done

Typical Poverty Scorecard

Indicator	Value	Points	Score
1. How many members does the household have?	A. Five or more	0	
	B. Four	6	
	C. Three	11	
	D. Two	17	
	E. One	20	
2. Do any household members ages 5 to 18 go to private school or private pre-school?	A. No	0	
	B. Yes	5	
	C. No members ages 5 to 18	7	
3. How many years of schooling has the female head/spouse completed?	A. Three or less	0	
	B. Four to eleven	2	
	C. Twelve or more	8	
	D. No female head/spouse	8	
4. How many household members work as employees with a written contract, as civil servants for the government, or in the military?	A. None	0	
	B. One	4	
	C. Two or more	13	
5. In their main occupation, how many household members are managers, administrators, professionals in the arts and sciences, mid-level technicians, or clerks?	A. None	0	
	B. One or more	8	
6. How many rooms does the residence have?	A. One to four	0	
	B. Five	2	
	C. Six	5	
	D. Seven	7	
	E. Eight or more	11	
7. How does the household dispose of sewage?	A. Ditch, other, or no bathroom	0	
	B. Simple hole, or directly into river, lake, or ocean	2	
	C. Septic tank not connected to public sewage/rainwater system	3	
	D. Septic tank connected to public sewage/rainwater system	4	
	E. Direct connection to public sewage/rainwater system	5	
8. Does the household have a refrigerator?	A. No	0	
	B. Yes, with one door	5	
	C. Yes, with two doors	10	
9. Does the household have a washing machine?	A. No	0	
	B. Yes	7	
10. Does the household have a cellular or land-line telephone?	A. None	0	
	B. Cellular but not land-line	5	
	C. Land-line but not cellular	6	
	D. Both	11	

\$2.50/Day/2005 PPP Poverty Line

PPI Score	Total Below the \$2.50/Day/2005 PPP Line	Total Above the \$2.50/Day/2005 PPP Line
0-4	81.8%	18.2%
5-9	77.8%	22.2%
10-14	66.1%	33.9%
15-19	49.0%	51.0%
20-24	37.2%	62.8%
25-29	23.9%	76.1%
30-34	15.4%	84.6%
35-39	8.6%	91.4%
40-44	5.2%	94.8%
45-49	3.2%	96.8%
50-54	2.1%	97.9%
55-59	1.2%	98.8%
60-64	1.2%	98.8%
65-69	0.4%	99.6%
70-74	0.6%	99.4%
75-79	0.0%	100.0%
80-84	0.0%	100.0%
85-89	0.0%	100.0%
90-94	0.0%	100.0%
95-100	0.0%	100.0%

Can we do better?

- This component of targeting is a pure prediction problem
- We fundamentally care about getting best predictive accuracy
- Let's use this example to illustrate the mechanics of prediction

Brazilian Data

- The data:
 - 44,787 data points
 - 53 variables
 - Not very wide?
- Median
 - Annual consumption (in dollars): 3918
 - 348.85 monthly income
- 6 percent below 1.90 poverty line
- 14 percent below the 3.10 poverty line

Coefficient on Number of Bedrooms

$\hat{\beta}$

0.4
0.3
0.2
0.1
0.0

0.00

0.25

0.50

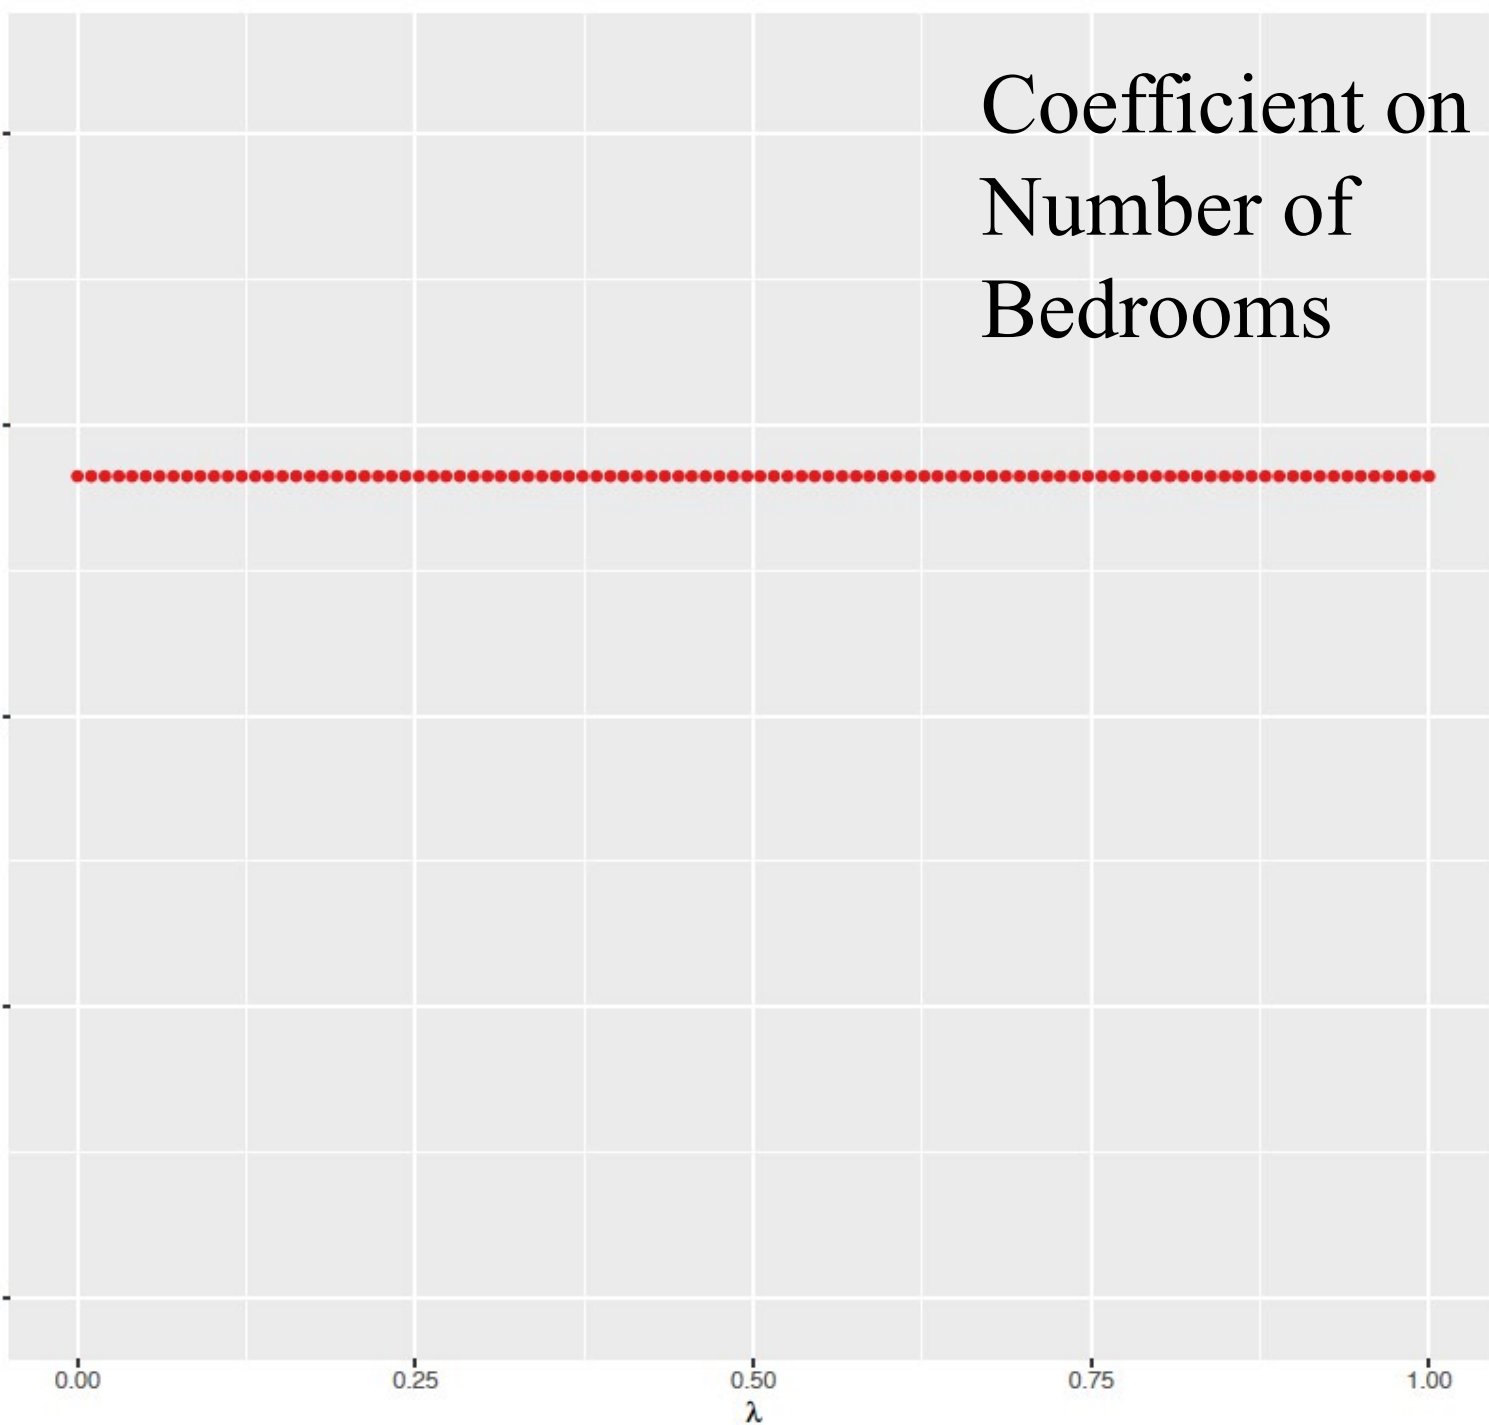
0.75

1.00


λ

method

• ols



Coefficient on Number of Bedrooms



Coefficient on Number of Bedrooms

What is this about?

