

# Classification

# Classification

- Examples
  - Who will default on their loans?
  - Identify a road or building in a picture.
  - Identify the party of a politician based on speech.
- Econometrics
  - Binary: OLS, logit, probit.
  - Multiple unordered classes: multinomial and conditional logit, multinomial probit.
- Machine Learning
  - Discriminant analysis
  - Naive Bayes
  - Support Vector Machines,...

# Classification

- Qualitative variables take values in an unordered set  $\mathcal{C}$ , such as:  
     $\text{eye color} \in \{\text{brown}, \text{blue}, \text{green}\}$   
     $\text{email} \in \{\text{spam}, \text{ham}\}.$
- Given a feature vector  $X$  and a qualitative response  $Y$  taking values in the set  $\mathcal{C}$ , the classification task is to build a function  $C(X)$  that takes as input the feature vector  $X$  and predicts its value for  $Y$ ; i.e.  $C(X) \in \mathcal{C}$ .
- Often we are more interested in estimating the *probabilities* that  $X$  belongs to each category in  $\mathcal{C}$ .

# Classification with large $p$

1. Regularized Logistic Regression
2. Support Vector Machine

# L1 regularized maximum likelihood

The standard maximum likelihood estimator solves

$$\max_{\beta} l(x, \beta).$$

Suppose that we have a linear index model, where  $l(x, \beta) = l(\beta'x)$ . The L1 regularized maximum likelihood solves

$$\max_{\beta} l(x, \beta) - \lambda \sum_{j=1}^p |\beta_j|,$$

where the intercept  $\beta_0$  is not penalized.

# Likelihood

For independent binary response data we can write the likelihood as

$$L(\theta, x) = F(x_1, \dots, x_n, \theta) = \prod_{i=1}^n F(x_i, \theta)$$

and the log likelihood as

$$\ln L(\theta, x) = l(\theta, x) = \sum_{i=1}^n \ln F(x_i, \theta).$$

- In the **Logit model**,  $F(z) = P[y = 1 \mid x]$  is the logistic function (let  $z = \beta'x$ )

$$F(z) = \frac{e^z}{1 + e^z}.$$

# L1 regularized logistic regression

The likelihood is

$$L = \prod_{i=1}^n (F[\beta'x_i])^{y_i} (1 - F[\beta'x_i])^{(1-y_i)}, \text{ so}$$

$$\begin{aligned} l &= \log L = \sum_{i=1}^n y_i \ln (F[\beta'x_i]) + (1 - y_i) \ln (1 - F[\beta'x_i]) \\ &= \sum_{i=1}^n \left( y_i \beta'x_i - \ln (1 + e^{\beta'x_i}) \right). \end{aligned}$$

and we solve




$$\max_{\beta} \sum_{i=1}^n \left( y_i \beta'x_i - \ln (1 + e^{\beta'x_i}) \right) - \lambda \sum_{j=1}^p |\beta_j|.$$

# Estimation

- Path algorithms such as LAR for lasso are more difficult.
- The R package *glmnet* can fit coefficient paths for very large logistic regression problems efficiently (large in  $N$  or  $p$ ). Their algorithms can exploit sparsity in the predictor matrix  $X$ , which allows for even larger problems.
- Stata 16 command *lasso* fits regularized logit (and probit and poisson) models.



# Example: Credit Card Default



## Credit Data

		<i>True Default Status</i>		
		No	Yes	Total
<i>Predicted Default Status</i>	No	9644	252	9896
	Yes	23	81	104
	Total	9667	333	10000

$(23 + 252)/10000$  errors — a 2.75% misclassification rate!

Some caveats:

- This is *training* error, and we may be overfitting. Not a big concern here since  $n = 10000$  and  $p = 2$ !

## Credit Data

		<i>True Default Status</i>		
		No	Yes	Total
<i>Predicted Default Status</i>	No	9644	252	9896
	Yes	23	81	104
	Total	9667	333	10000

$(23 + 252)/10000$  errors — a 2.75% misclassification rate!

Some caveats:

- This is *training* error, and we may be overfitting. Not a big concern here since  $n = 10000$  and  $p = 2$ !
- If we classified to the prior — always to class **No** in this case — we would make  $333/10000$  errors, or only 3.33%.

## Credit Data

		<i>True Default Status</i>		
		No	Yes	Total
<i>Predicted Default Status</i>	No	9644	252	9896
	Yes	23	81	104
	Total	9667	333	10000

$(23 + 252)/10000$  errors — a 2.75% misclassification rate!

Some caveats:

- This is *training* error, and we may be overfitting. Not a big concern here since  $n = 10000$  and  $p = 2$ !
- If we classified to the prior — always to class **No** in this case — we would make  $333/10000$  errors, or only 3.33%.
- Of the true **No**'s, we make  $23/9667 = 0.2\%$  errors; of the true **Yes**'s, we make  $252/333 = 75.7\%$  errors!

## Types of errors

**False positive rate:** The fraction of negative examples that are classified as positive — 0.2% in example.

**False negative rate:** The fraction of positive examples that are classified as negative — 75.7% in example.

We produced this table by classifying to class **Yes** if


$$\widehat{\Pr}(\text{Default} = \text{Yes} | \text{Balance}, \text{Student}) \geq 0.5$$

We can change the two error rates by changing the threshold from 0.5 to some other value in  $[0, 1]$ :

$$\widehat{\Pr}(\text{Default} = \text{Yes} | \text{Balance}, \text{Student}) \geq \textit{threshold},$$


and vary *threshold*.

## Varying the *threshold*




In order to reduce the false negative rate, we may want to reduce the threshold to 0.1 or less.

ROC Curve



The *ROC plot* displays both simultaneously.

ROC Curve



The *ROC plot* displays both simultaneously.

Sometimes we use the *AUC* or *area under the curve* to summarize the overall performance. Higher *AUC* is good.



# Confusion matrix

- Table of true and predicted classes.
- Terminology

<i>Predicted class</i>		True class		
		No	Yes	Total
	No	True Neg: (TN)	False Neg. (FN)	N*
	Yes	False Pos. (FP)	True Pos. (TP)	P*
	Total	N	P	

Name	Definition	Synonyms
False Pos. rate	FP/N	Type I error, 1–Specificity
True Pos. rate	TP/P	1–Type II error, power, sensitivity, recall
Pos. Pred. value	TP/P*	Precision, 1–false discovery proportion
Neg. Pred. value	TN/N*	

# Support Vector Machines

Here we approach the two-class classification problem in a direct way:

*We try and find a plane that separates the classes in feature space.*

If we cannot, we get creative in two ways:

- We soften what we mean by “separates”, and
- We enrich and enlarge the feature space so that separation is possible.


# What is a Hyperplane?

- A hyperplane in  $p$  dimensions is a flat affine subspace of dimension  $p - 1$ .
- In general the equation for a hyperplane has the form


$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p = 0$$

- In  $p = 2$  dimensions a hyperplane is a line.
- If  $\beta_0 = 0$ , the hyperplane goes through the origin, otherwise not.
- The vector  $\beta = (\beta_1, \beta_2, \dots, \beta_p)$  is called the normal vector — it points in a direction orthogonal to the surface of a hyperplane.

# Hyperplane in 2 Dimensions




## Separating Hyperplanes



- If  $f(X) = \beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p$ , then  $f(X) > 0$  for points on one side of the hyperplane, and  $f(X) < 0$  for points on the other.
- If we code the colored points as  $Y_i = +1$  for blue, say, and  $Y_i = -1$  for mauve, then if  $Y_i \cdot f(X_i) > 0$  for all  $i$ ,  $f(X) = 0$  defines a *separating hyperplane*.

# Maximal Margin Classifier

Among all separating hyperplanes, find the one that makes the biggest gap or margin between the two classes.



Constrained optimization problem

$$\text{maximize } M$$

$$\beta_0, \beta_1, \dots, \beta_p$$

$$\text{subject to } \sum_{j=1}^p \beta_j^2 = 1,$$


$$y_i(\beta_0 + \beta_1 x_{i1} + \dots + \beta_p x_{ip}) \geq M$$

for all  $i = 1, \dots, N$ .



This can be rephrased as a convex quadratic program, and solved efficiently. The function `svm()` in package `e1071` solves this problem efficiently

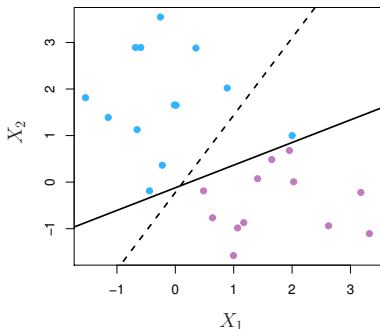
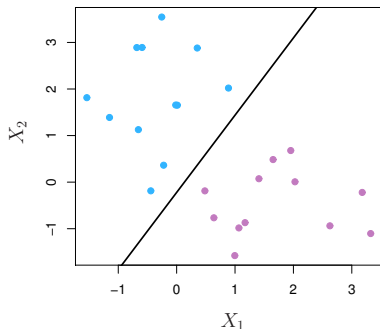
## Non-separable Data



The data on the left are not separable by a linear boundary.

This is often the case, unless  $N < p$ .

## Noisy Data





Sometimes the data are separable, but noisy. This can lead to a poor solution for the maximal-margin classifier.

The *support vector classifier* maximizes a *soft* margin.



# Support Vector Classifier

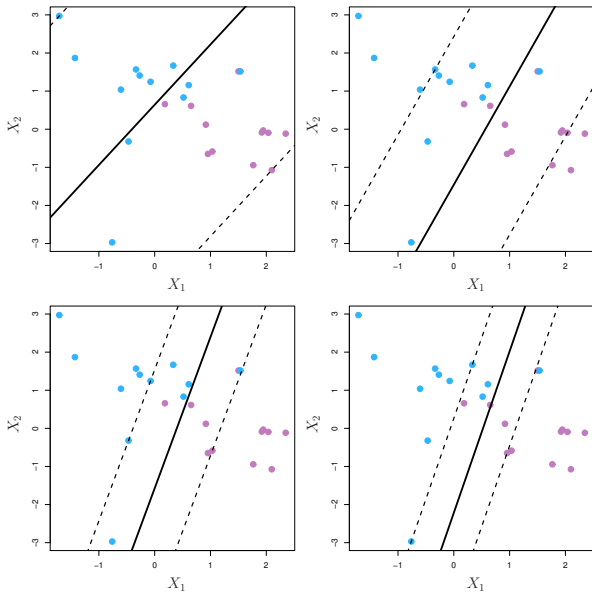


$$\underset{\beta_0, \beta_1, \dots, \beta_p, \epsilon_1, \dots, \epsilon_n}{\text{maximize}} \quad M \quad \text{subject to} \quad \sum_{j=1}^p \beta_j^2 = 1,$$

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i),$$

$$\epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq \textcolor{red}{C}, \quad \text{NB: different } C \text{ than below.}$$

$C$  is a regularization parameter



# Separating hyperplanes

A hyperplane or affine set  $L$  is defined by the equation

$$f(x) = \beta_0 + x' \beta = 0.$$

If we are in  $\mathbb{R}^2$ , this is a line. Some properties

1. For any two points  $x_1$  and  $x_2$  lying in  $L$ ,  $v = x_1 - x_2$  is parallel to the hyperplane. Since  $(x_1 - x_2)' \beta = 0$ ,  $\beta^* = \beta / \|\beta\|$  is the vector normal to the surface of  $L$ .
2. The signed distance of any point  $x$  to  $L$  is given by  $(x' \beta - \beta_0) / \|\beta\|$

$$(x - x_1)' \beta^* = \frac{1}{\|\beta\|} (x - x_1)' \beta = \frac{1}{\|\beta\|} (x' \beta + \beta_0) = \frac{1}{\|\beta\|} f(x).$$

Hence  $f(x)$  is proportional to the signed distance from  $x$  to the hyperplane defined by  $f(x) = 0$ .

Our training data consists of  $n$  pairs  $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$ , with  $x_i \in \mathbb{R}^p$  and  $y_i \in \{-1, 1\}$ . Consider the optimization problem of choosing the maximal separating margin

$$\max_{\beta, \beta_0} M$$

subject to

$$y_i (x_i' \beta + \beta_0) \frac{1}{\|\beta\|} \geq M, \quad \forall i,$$

or equivalently,

$$y_i (x_i' \beta + \beta_0) \geq M \|\beta\|, \quad \forall i.$$

We can arbitrarily set the scaling of  $\|\beta\|$  to  $1/M$ . Thus we get ( $\operatorname{argmin} \|\beta\| = \operatorname{argmin} \frac{1}{2} \|\beta\|^2$ )

$$\min_{\beta, \beta_0} \frac{1}{2} \|\beta\|^2,$$

subject to

$$y_i (x_i' \beta + \beta_0) \geq 1, \quad i = 1, 2, \dots, n.$$

( $\|\beta\| = 1$  m better than  $\|\beta\| = 1$  km).

The Lagrangeian is

$$\min_{\beta, \beta_0} \mathcal{L} = \frac{1}{2} \|\beta\|^2 - \sum_{i=1}^n \alpha_i [y_i (x_i' \beta + \beta_0) - 1].$$

The first order conditions w.r.t  $\beta$  and  $\beta_0$  are

$$\begin{aligned} \beta &= \sum_{i=1}^n \alpha_i y_i x_i, \\ 0 &= \sum_{i=1}^n \alpha_i y_i. \end{aligned} \tag{1}$$

Substituting in  $\mathcal{L}$  we obtain the so-called Wolfe dual

$$\mathcal{L}_D = \underbrace{\frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j x_i' x_j}_{\beta' \beta} - \sum_{i=1}^n \alpha_i y_i \underbrace{\sum_{j=1}^n \alpha_j y_j x_j' x_i}_{\beta'} - \beta_0 \underbrace{\sum_{i=1}^n \alpha_i y_i}_{=0} + \sum_{i=1}^n \alpha_i$$

$$\min_{\alpha} \mathcal{L}_D = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{k=1}^n \alpha_i \alpha_k y_i y_k x_i' x_k. \quad (2)$$

subject to

$$\alpha_i \geq 0 \text{ and } \sum_{i=1}^n \alpha_i y_i = 0.$$

This is a simpler convex optimization problem, for which standard software can be used. In addition the solution must satisfy the Kuhn–Tucker conditions, which include the above conditions and

$$\alpha_i [y_i (x_i' \beta + \beta_0) - 1] = 0, \quad i = 1, 2, \dots, n.$$



Thus

- if  $\alpha_i > 0$ , then  $y_i (x_i' \beta + \beta_0) = 1$ , in other words,  $x_i$  is on the boundary,
- if  $y_i (x_i' \beta + \beta_0) > 1$  then  $\alpha_i = 0$  and  $x_i$  is not on the boundary.
- From equation 1, we see that the solution vector is defined in terms of a linear combination of the support points  $x_i$  — those points defined to be on the boundary with  $\alpha_i > 0$ .

# Support Vector Machines

Suppose now that the classes overlap in feature space. One way to deal with the overlap is to still maximize  $M$ , but allow for some points to be on the wrong side of the margin. Define the slack variables

$$\xi = (\xi_1, \xi_2, \dots, \xi_n).$$



We now solve

$$\min_{\beta, \beta_0} \frac{1}{2} \|\beta\|^2,$$

subject to

$$y_i (x_i' \beta + \beta_0) \geq 1 - \xi_i, \quad \forall i,$$

$$\xi_i \geq 0, \quad \sum_{i=1}^n \xi_i \leq \text{constant},$$

resulting in Lagrangeian

$$\mathcal{L} = \frac{1}{2} \|\beta\|^2 + C \sum_{i=1}^n \xi_i - \sum_{i=1}^n \alpha_i [y_i (x_i' \beta + \beta_0) - (1 - \xi_i)] - \sum_{i=1}^n \mu_i \xi_i.$$



The first-order conditions w.r.t.  $\beta$  and  $\beta_0$  are as in equation (1) and the new (wrt  $\xi_i$ )

$$\alpha_i = C - \mu_i, \quad \forall i,$$

as well as  $\alpha_i, \mu_i, \xi_i > 0 \forall i$ . The dual Lagrangeian is as in equation (2). We minimize this subject to

$$\begin{aligned} 0 &\leq \alpha_i \leq C \\ 0 &= \sum_{i=1}^n \alpha_i y_i. \end{aligned}$$

The Kuhn-Tucker conditions include

$$\begin{aligned} \alpha_i [y_i (x_i' \beta + \beta_0) - (1 - \xi_i)] &= 0, \\ \mu_i \xi_i &= 0, \\ y_i (x_i' \beta + \beta_0) - (1 - \xi_i) &\geq 0, \end{aligned}$$

for all  $i$ .

- The coefficient vector again has the form

$$\hat{\beta} = \sum_{i=1}^n \hat{\alpha}_i y_i x_i$$

with nonzero coefficients  $\hat{\alpha}_i$  only for those observations  $i$  for which  $y_i (x_i' \beta + \beta_0) = (1 - \xi_i)$  due to the first Kuhn-Tucker condition.


- These observations are called the support vectors, since  $\hat{\beta}$  is represented in terms of them alone. Among these support points, some will lie on the edge of the margin ( $\xi_i = 0$ ), and hence will be characterized by  $0 \leq \hat{\alpha}_i \leq C$  ( $\hat{\alpha}_i = C - \hat{\mu}_i$ ). The remainder will have  $\alpha_i = C$ .
- $C$  is a tuning parameter. A large value of  $C$  will discourage any positive  $\xi_i$ .

- Once you have the  $\hat{\alpha}_i$ , classification is easy

$$\begin{aligned} f(x) &= \beta_0 + x'\beta = \beta_0 + x' \sum_{i=1}^n \hat{\alpha}_i y_i x_i \\ &= \beta_0 + \sum_{i=1}^n \hat{\alpha}_i y_i x' x_i. \end{aligned}$$

- Only observations that are support vectors ( $\hat{\alpha}_i > 0$ ) matter.
- Positive correlation ( $x'x_i > 0$ ) with positive examples ( $y_i = 1$ ) and negative correlation ( $x'x_i < 0$ ) with negative examples ( $y_i = -1$ ) increase  $f(x)$  chance of being coded as positive
- Correlation with examples  $K(x, x_i)$  can be made more general, see below.

## Linear boundary can fail




Sometime a linear boundary simply won't work, no matter what value of  $C$ .

The example on the left is such a case.

What to do?

# Radial Kernel

$$K(x_i, x_{i'}) = \exp(-\gamma \sum_{j=1}^p (x_{ij} - x_{i'j})^2).$$





$$f(x) = \beta_0 + \sum_{i \in \mathcal{S}} \hat{\alpha}_i K(x, x_i)$$

Implicit feature space;  
very high dimensional.



Controls variance by  
squashing down most  
dimensions severely

## Example: Heart Data



ROC curve is obtained by changing the threshold 0 to threshold  $t$  in  $\hat{f}(X) > t$ , and recording *false positive* and *true positive* rates as  $t$  varies. Here we see ROC curves on training data.

## Example continued: Heart Test Data



## SVMs: more than 2 classes?

The SVM as defined works for  $K = 2$  classes. What do we do if we have  $K > 2$  classes?

**OVA** One versus All. Fit  $K$  different 2-class SVM classifiers  $\hat{f}_k(x)$ ,  $k = 1, \dots, K$ ; each class versus the rest. Classify  $x^*$  to the class for which  $\hat{f}_k(x^*)$  is largest.

**OVO** One versus One. Fit all  $\binom{K}{2}$  pairwise classifiers  $\hat{f}_{k\ell}(x)$ . Classify  $x^*$  to the class that wins the most pairwise competitions.

Which to choose? If  $K$  is not too large, use OVO.



## Which to use: SVM or Logistic Regression

- When classes are (nearly) separable, SVM does better than LR. So does LDA.
- When not, LR (with ridge penalty) and SVM very similar.
- If you wish to estimate probabilities, LR is the choice.
- For nonlinear boundaries, kernel SVMs are popular. Can use kernels with LR and LDA as well, but computations are more expensive.

# Readings

## Classification:

*Introduction to Statistical Learning*: Chapter 4.

## Support Vector Machines:

*Introduction to Statistical Learning*: Chapter 9.

*Elements of Statistical Learning*: Chapter 12